

Task: Fix Missing Experience Data for 1,000 Candidates

Data Engineering - TASK

Context

We have a CSV, `missing_duration_candidates.csv`, with 1,000 candidates whose experience data is broken. 972 have an experience entry but no usable duration. 28 have no experience at all likely a failed scrape. Both groups need to be fixed.

What to do

1. Study the schema. Hit the two engine API URLs below and look at the response. That's the shape your final data needs to match.
2. For each candidate, scrape their `linkedin_url` with the Apify actor linked below.
3. Parse the result into the fields listed under Data Format, and write it back to the DB under the same candidate ID (don't create new rows).
4. Log every candidate as a success or failure, with a reason for failures. Don't let failures fail silently.

Data format

Every position under experience must have these fields. This is what 'fixed' means:

Field	Notes
<code>role</code> , <code>company</code>	Required, non-empty.
<code>start_date</code>	Month/year is fine, day can default to the 1st.
<code>end_date</code>	null if the role is current — don't invent a fake end date.
<code>is_current</code>	true/false.
<code>duration_months</code> / <code>duration_text</code>	Computed from <code>start_date</code> and <code>end_date</code> — not hand-typed.

Watch out for

- Current roles: In the two reference profiles, the engine itself computes `duration_months` inconsistently for ongoing roles (one shows the correct month count, others show 0). Don't copy that pick one rule (e.g. months from `start_date` to today) and apply it everywhere.
- The 28 `empty_experience` candidates check whether LinkedIn genuinely shows no experience for them, or whether it's a stale/failed scrape worth retrying, before writing an empty array.
- Some companies won't have a resolvable LinkedIn company page; that's fine, just don't crash on it.

Failures

Expect some LinkedIn URLs to fail (private, dead, renamed, scraper timeout). Aim for at least 950/1,000 fixed. Anything below that, or a failure rate above 10%, stop and flag it rather than pushing through that, usually means something's wrong with the run itself.

MANDATORY

Build GET `/api/saral/profile/{id}` that returns the same response shape as the engine endpoint, sourced from our DB.

Check it by comparing your response to the engine's for the same candidate id.

Done when

- 950+ candidates have valid experience in the DB (role, company, dates, duration) per the format above.
- Failed candidates are logged with a reason.
- Code is in a repo with a short README on how to run it.

References

- Apify actor: console.apify.com/actors/LpVuK3Zozwuipa5bp
- Engine API example 1:
engine.saralhire.ai/api/saral/profile/48125f49-d89c-407b-8b38-8dec96c9ab99
- Engine API example 2:
engine.saralhire.ai/api/saral/profile/b432fd65-121b-4af7-a0de-1b625826c7c2
- Source CSV: `missing_duration_candidates.csv`
- DB credentials: provided separately